

Neighbor-Sensitive Hashing

Yongjoo Park Michael Cafarella Barzan Mozafari

University of Michigan, Ann Arbor

Nov 05, 2015

k -Nearest Neighbors Search (k NN)



Database



k -Nearest Neighbors Search (k NN)

$$q = \{120, 65, \dots, 212\}$$



Database



k -Nearest Neighbors Search (k NN)

$$q = \{120, 65, \dots, 212\}$$




Database



$$v1 = \{\dots\}$$




$$v2 = \{\dots\}$$




$$v3 = \{\dots\}$$



$$v4 = \{\dots\}$$




$$v5 = \{\dots\}$$




$$v6 = \{\dots\}$$

k -Nearest Neighbors Search (k NN)

$$q = \{120, 65, \dots, 212\}$$




Database



$$v1 = \{\dots\}$$




$$v2 = \{\dots\}$$




$$v3 = \{\dots\}$$




$$v4 = \{\dots\}$$



$$v5 = \{\dots\}$$



$$v6 = \{\dots\}$$




k -Nearest Neighbors Search (k NN)


$$q = \{120, 65, \dots, 212\}$$




Database



$$v1 = \{\dots\}$$



$$v2 = \{\dots\}$$




$$v3 = \{\dots\}$$



$$v4 = \{\dots\}$$



$$v5 = \{\dots\}$$




$$v6 = \{\dots\}$$



Exact k NN for high-dimensional vectors is **Slow**.

- (Kashyap *KDD*'11) took more than 1 min for 10M objects.



Hashing-based k NN



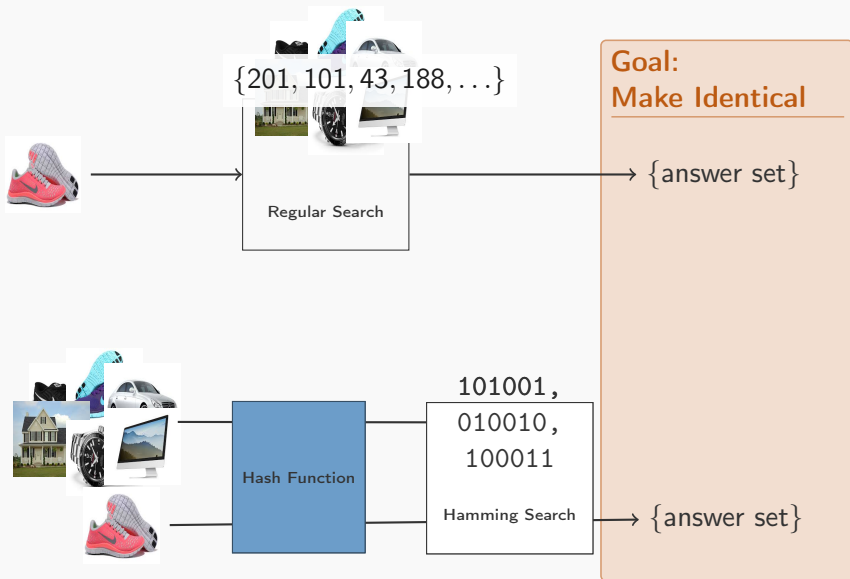
Hashing-based k NN



Hashing-based k NN



Hashing-based k NN



Previous Approaches

- 2004, Locality Sensitive Hashing (LSH), Datar et al.

Previous Approaches

- 2004, Locality Sensitive Hashing (LSH), Datar et al.
- 2009, Spectral Hashing (SH), Weiss et al.

Previous Approaches

- 2004, Locality Sensitive Hashing (LSH), Datar et al.
- 2009, Spectral Hashing (SH), Weiss et al.
- 2011, Anchor Graph Hashing (AGH), Liu et al.

Previous Approaches

- 2004, Locality Sensitive Hashing (LSH), Datar et al.
- 2009, Spectral Hashing (SH), Weiss et al.
- 2011, Anchor Graph Hashing (AGH), Liu et al.
- 2012, Spherical Hashing (SpH), Heo et al.
Kernelized Supervised Hashing (KSH), Liu et al.

Previous Approaches

- 2004, Locality Sensitive Hashing (LSH), Datar et al.
- 2009, Spectral Hashing (SH), Weiss et al.
- 2011, Anchor Graph Hashing (AGH), Liu et al.
- 2012, Spherical Hashing (SpH), Heo et al.
Kernelized Supervised Hashing (KSH), Liu et al.
- 2013, Compressed Hashing (CH), Lin et al.
Complementary Projection Hashing (CPH), Jin et al.


Previous Approaches

- 2004, Locality Sensitive Hashing (LSH), Datar et al.
- 2009, Spectral Hashing (SH), Weiss et al.
- 2011, Anchor Graph Hashing (AGH), Liu et al.
- 2012, Spherical Hashing (SpH), Heo et al.
Kernelized Supervised Hashing (KSH), Liu et al.
- 2013, Compressed Hashing (CH), Lin et al.
Complementary Projection Hashing (CPH), Jin et al.
- 2014, Data Sensitive Hashing (DSH), Jagadish et al.

Previous Approaches


- 2004, Locality Sensitive Hashing (LSH), Datar et al.
- 2009, Spectral Hashing (SH), Weiss et al.
- 2011, Anchor Graph Hashing (AGH), Liu et al.
- 2012, Spherical Hashing (SpH), Heo et al.
Kernelized Supervised Hashing (KSH), Liu et al.
- 2013, Compressed Hashing (CH), Lin et al.
Complementary Projection Hashing (CPH), Jin et al.
- 2014, Data Sensitive Hashing (DSH), Jagadish et al.

LSH




Previous Approaches

- 2004, Locality Sensitive Hashing (LSH), Datar et al.
- 2009, Spectral Hashing (SH), Weiss et al.
- 2011, Anchor Graph Hashing (AGH), Liu et al.
- 2012, Spherical Hashing (SpH), Heo et al.
Kernelized Supervised Hashing (KSH), Liu et al.
- 2013, Compressed Hashing (CH), Lin et al.
Complementary Projection Hashing (CPH), Jin et al.
- 2014, Data Sensitive Hashing (DSH), Jagadish et al.



Previous Approaches

- 2004, Locality Sensitive Hashing (LSH), Datar et al.
- 2009, Spectral Hashing (SH), Weiss et al.
- 2011, Anchor Graph Hashing (AGH), Liu et al.
- 2012, Spherical Hashing (SpH), Heo et al.
Kernelized Supervised Hashing (KSH), Liu et al.
- 2013, Compressed Hashing (CH), Lin et al.
Complementary Projection Hashing (CPH), Jin et al.
- 2014, Data Sensitive Hashing (DSH), Jagadish et al.




Previous Approaches

- 2004, Locality Sensitive Hashing (LSH), Datar et al.

Good idea for:

- ▶ Sorting all data items

- 2012, Spherical Hashing (SpH), Heo et al.
Kernelized Supervised Hashing (KSH), Liu et al.
- 2013, Compressed Hashing (CH), Lin et al.
Complementary Projection Hashing (CPH), Jin et al.
- 2014, Data Sensitive Hashing (DSH), Jagadish et al.



Previous Approaches

- 2004, Locality Sensitive Hashing (LSH), Datar et al.

Good idea for:


- ▶ Sorting all data items

- 2012, Spherical Hashing (SpH), Heo et al.



Remember:

- ▶ We are interested in only k items.



- 2014, Data Sensitive Hashing (DSH), Jagadish et al.



Our Approach




Our Approach




More Separators between Neighbors

- Neighbors close to the query are better distinguished
- Low Resolution for distant items \implies **No Problem!**

Algorithmic Details



(a) Regular Hashing




(b) Hashing with NST

A special Non-Linear Transformation

- Expand the distance between Neighbors
- Effectively, more hash functions for Neighbors

Recall Improvement over State-of-the-art




Hashcode size vs. Recall Improvement

- ▶ Up to 10% recall improvement over SpH [2]
- ▶ Up to 15% recall improvement over LSH [1]

(The SIFT dataset)

Speed Improvement over State-of-the-art



Target Recall vs. Time Reduction

- ▶ Up to 22% time reduction over CPH [3]
- ▶ Up to **67%** time reduction over LSH [1]

(The SIFT dataset)

Questions

- [1] **M. Datar, N. Immorlica, P. Indyk, and V. S. Mirrokni.**
Locality-sensitive hashing scheme based on p-stable distributions.
In *SoCG*, 2004.
- [2] **J.-P. Heo, Y. Lee, J. He, S.-F. Chang, and S.-E. Yoon.**
Spherical hashing.
In *CVPR*, 2012.
- [3] **Z. Jin, Y. Hu, Y. Lin, D. Zhang, S. Lin, D. Cai, and X. Li.**
Complementary projection hashing.
In *ICCV*, 2013.